

Report “Distributed SLAM”

Bertelli Davide, Hueller Jhonny

Abstract— This report presents the results, solutions and implementation of a distributed SLAM system on a natural environment. The data were acquired through terrestrial agents and processed on-board. The agents were allowed to communicate within a certain range. Each agent had at its disposal a Kalmann Filter for data denoising and an Extendend Kalmann Filter for position estimation. The results shows a good accuracy in reconstructing the environment both in a 3D point cloud and as an occupancy map. The estimation of the location resulted unsatisfactory.

I. INTRODUCTION

The SLAM, acronym for Simultaneous Localisation And Mapping, is a method employed on autonomous vehicles which allows to gather data from the environment while localising the agent at the same time. Furthermore, it allows to the agent to map unknown environments and update its knowledge of the surroundings. The map resulting will then be used for further steps in the autonomous navigation like path planning. Agents exploiting the SLAM algorithm have been vastly employed in several working realities which space from the industrial up to civil solutions. The Distributed SLAM is a variation of the SLAM method, in which multiple agents contribute in mapping chunks of the same area to have a faster and more precise result than the single agent case. This method aims at exploiting at maximum the capabilities of the different agents involved allowing to explore areas which are unexplorable otherwise.

A. Problem Formulation

The problem tried to solve in this report regards the development of a Distributed SLAM algorithm to be deployed in natural environments. Specifically the environment was assumed to be a forest and, by definition, that the surroundings contain several elements with different but recurrent shapes, like the trees. Furthermore it was kept into account the possible presence of pits which could doom the autonomous navigation of the agents, by posing them into an unrecoverable location. The agents employed are assumed to be terrestrial and that their initial locations are known a priori. The objective is to explore as long as possible the surroundings in order to map them. The output of this process will then be a 3D cloud of points representing the mapped environment and a 2D occupancy map representing the elements detected, their locations and encumbrance. Lastly the environment it is supposed to have areas in which the GPS signal cannot be received.

II. ADOPTED MODELS

A. Communication System

The Communication System adopted is a Robot Network in which each agent can transmit its data to any nearby agent within a specified communication range. There is no master among the devices, hence a fully distributed network is achieved. Indeed each agent acquires the data from its sensors, process them onboard and then transmits them only if any other agent has been detected within range. The informations transmitted consists of the point cloud representing the environment up to the time of transmission and the occupancy map in the local coordinate system together with its history of poses.

B. System Model

The robot used in the development of the solution and during the simulations is a terrestrial one. It has been modelled just for simulation purposes and so its shape is simple even though it has the same dimensions of a common RC car, as shown in Figure 1. The sensors onboard of the agent are supposed to be, a 3D Lidar, an accelerometer and a gyroscope. Thanks to these sensors it will be possible to map the environment around and know the linear and angular velocities of the agent together with the rotation on its axis. For clarity, the robot has been simulated moving over the XY plane while rotating around the Z axis. The Lidar sensor in use aims at simulating the characteristics of the RS-LIDAR-16 sensor, but due to computational demands during simulations its horizontal resolution has been reduced to allow reasonable computation times. The simulated Lidar sensor has the beam origin in the center of the upper part of the agent at an height of 0.135[m] and has the following features:

Range	15 meters
Horizontal Beams	360
Horizontal FoV	$[-3.14, 3.14]$ radians
Vertical Beams	16
Vertical FoV	$[-0.26, 0.26]$ radians




Figure 1: Schema of the simulated agent.

The actuators onboard of the device are supposed to be four electrical motors posed nearby each wheel and providing traction both in the frontal and rear axles. Furthermore, only the front wheels are supposed to rotate over the Z axis to turn the agent in the desired direction.

III. SOLUTION

Our proposed solution consists in a distributed network of robots in which each agent will acquire data from the environment and process them autonomously. Communication is supported only for agent within a certain distance known as a communication range. The pipeline that each agent will execute during its lifespan can be summarized in the following steps:

1. Acquire data from the 3D Lidar sensor.
2. Remove sensor noise from the data.
3. Save cleaned point cloud and generate 2D occupancy map.
4. Check if there are agents nearby. If it is the case:
 - (a) Acquire data from nearby agents.
 - (b) Update local point cloud.
 - (c) Update occupancy map.
5. Sample next location to explore using the occupancy map.
6. Execute the maneuvers to reach the location. In this step move using the GPS location if available, otherwise estimate the location during the movement.
7. Once arrived update the current location and repeat.

In developing the solution to the Distributed SLAM in a forest-like environment we had to tackle several issues, in particular the main ones were:

- Motion Model of the agent.
- Removing Lidar Sensor noise.
- Roadmap Generation and Path Planning.
- Estimating agent location in an area without GPS coverage.
- Possible discontinuous terrain.
- Communication among agents.

A. Motion Model

Regarding the motion model, due to the simplicity of the shape and the inputs at our disposal, we decided to employ the Differential Drive kinematic model which represented our robot as a two-wheeled one rotated of a quantity θ with respect to the X axis and revolving around the Z one as shown in Figure 2. Furthermore, it allowed to control our agents motion by providing the desired linear and angular velocities other than specifying the wheels radius and axle separation.




Figure 2: Differential Drive example.

B. Lidar Noise Removal

Due to the data acquisition from a sensor it was not possible to omit the possible presence of noise in the informations retrieved. For this reason we assumed the presence of an additive gaussian noise in the input stream and to counteract its effect we decided to use a linear Kalman Filter to estimate the true value of the data. Specifically, the equations ruling the Kalman Filter in use are the following.

$$\hat{x} = A\hat{x} + Bu \quad (1)$$

$$P = APA^{-1} + Q \quad (2)$$

$$K = PH^T(HPH^T + R)^{-1} \quad (3)$$

$$\hat{x} = \hat{x} + K(Z - H\hat{x}) \quad (4)$$

$$P = P - KHP \quad (5)$$

Where \hat{x} is the estimated state, A is the state transition matrix, B is the input matrix, u is the input vector, Q is the process noise covariance, K is the Kalman gain, H is the observation matrix, R is the observation covariance and Z is the measurement taken from the sensor.

C. Roadmap and Path Planning

Concerning the Roadmap Generation and Path Planning issues, we decided to make the agent chose its next location through a uniform random sample operation over the free space areas within a maximum distance equal to the lidar range. In this way the worst scenario would still have the agent in a free space area. Furthermore, we also supposed that the GPS location could incur into some errors and so we decided to allow the agent to approximate its location to the nearest free one through gaussian sampling in case the initial location would result into an occupied or invalid one. Once the initial and arrival locations are in a free space area, the agent will then generate the roadmap necessary to reach the arrival position. The algorithm chosen to carry out this task is the RTT*. The choice of this algorithm has been lead due to its ability to find paths even through narrow apertures between detected objects and also due to its inner nature of being a tree graph which will depart from the initial location, exploring and re-wiring itself until it reaches its goal. This latter behaviour has been decisive in the choice, especially against the common PRM which will randomly sample all the environment generating a bunch of useless data. Lastly, to ensure a certain fluidity in the agent motion we constrained the sample of the next location using spatial boundaries like the ranges of arrival rotation and the minimum and maximum x-y values.

D. Agent Position Estimation

Regarding the position of the agents in use we are exploiting their absolute location in the space and we keep it updated thanks to the data retrieved by the GPS onboard. Unfortunately in natural environment it is not always possible to rely on these measurements, so we needed also a way to estimate the location of the robots relying only on the measurements retrieved and the motion model of the robot itself. To achieve this task we decided to employ an Extended Kalmann Filter which is specific for non-linear models. Its strength is due to the linearization through derivation of the non-linear functions ruling the measurements acquisition and the states transition. The equations describing its behaviour are the same of the linear Kalmann Filter, with the only difference that the covariance of the state and the observation are computed deriving the respective functions.

E. Discontinuous Terrain

Working on a natural environment must also imply that the surface on which the agents will move cannot be homogeneous and so the presence of pits or holes must be taken into account. For this reason in the processing of the 3D point cloud coming from the Lidar we appended some lines of code which will perform an automatic detection of the points representing the ground and then a further classification of these points into the pits ones. The ground detection has been carried out through the SMRF algorithm, while the points representing the pits have been classified depending by their Z distance from the origin of the Lidar beams. Doing so we can ensure that the agents will avoid such detected pits by embedding their points into the occupancy map generated by the original point cloud, tricking the robots into observing the pits as occupied space.

F. Communication

Being our system distributed over a multitude of agents it is also necessary to establish some procedure of communication. In particular we have assumed that two arbitrary agents can share informations if they are within a certain communication range. To simulate this behaviour we supposed that the agents have onboard some kind of sensor which act as a proximity one telling at which distance and orientation there is an agent. These informations have been practically provided to the robots through a third kind of agent which is omniscient regarding the global positioning of the other agents and that can access their inner class data. It is important to point out that this is a mere program and not a simulated sensor itself. Thanks to its presence it will then be possible to virtually share the data between the existing agents.

IV. IMPLEMENTATION DETAILS

In this section of the report will be presented how the system has been simulated and implemented, together with its architecture.

A. Software In Use

For the development of our solution we used a local machine running Matlab R2021b for creating and executing

the scripts ruling the agents, performing the data processing and allowing the connection to the simulation environment. This latter one is generated by the simulation software Gazebo v11.11 which provided the physical constraints of the objects and environment dynamics. The worlds over which the simulation runs have been manually created using Blender v3.2. Furthermore the bidirectional communication between Matlab and Gazebo has been possible using the ROS framework integrated with Gazebo, specifically ROS Noetic v1.15.14. The software architecture is summarized and proposed in Figure 3.

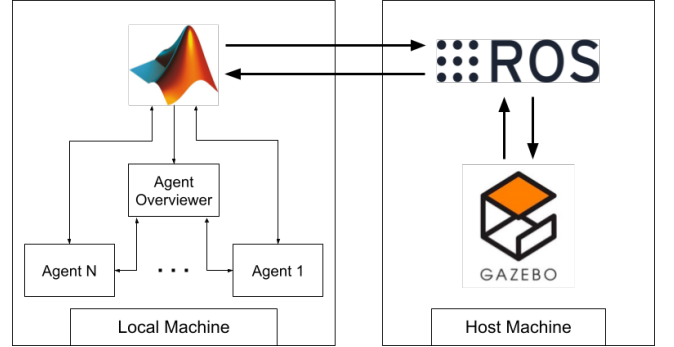


Figure 3: Schema of the Simulation Architecture

B. Software Components

As Shown in Figure 3 matlab rules over two classes of objects: the agents and the agentOverview. The agent class is the one responsible for the data acquisition and processing, roadmap generation, path planning and maneuver execution. It communicates with the simulated environment through Matlab using the ROS framework to retrieve sensible data from gazebo. The data of interest are the simulated 3D Lidar output and the odometry of the agent. The other class ruled by the Matlab environment is the agentOverview which accesses all the properties of the agents existing and exploits their knowledge and methods to share data with all the other agents which satisfies the communication criterion. It is important to notice how this class do not interfaces directly to the simulation through Matlab, but delegates the agents into taking the data it needs from the simulated environment. This behaviour is used to simulate a digital proximity sensor supposed onboard of each agent. Lastly, this class also embeds the informations regarding which areas of the global map do not allow for GPS data retrieval and forces the agents to estimate their locations.

V. RESULTS

In this section of the report will be shown the results achieved by our solution. For sake of clarity we will present results regarding the different steps of the workflow executed by each agent and lastly the global results. We will also put focus on the issues tackled in Section III.. Specifically, we will start discussing about the 3D Lidar data denoising, then about the pits detection, later on the sharing data among the agents, thereafter about the roadmap generation and lastly about the estimation of the agent location in areas without GPS.

A. Lidar Data Denoising

In processing the incoming 3D point cloud from the Lidar sensor we decided to estimate its true values against the noise exploiting a Linear Kalmann filter, but before testing it directly against the Lidar data we wanted to have a clear understanding of the goodness of our implementation. For this reason we generated samples from a sine function to which we have summed a gaussian noise with 0 mean and 0.1 variance. On these samples we applied the Kalmann Filter and retained its estimates. The result can be seen in Figure 4 and Figure 5. In the first figure we can observe how the estimated values in output from the Kalmann Filter still exhibit the noise behaviour, but correctly follows the curves of the sine wave. In particular by increasing the values of the observation covariance matrix in the Kalmann Filter it is possible to achieve a smoother behaviour of the estimates, but after several trials we decided that its current value, 0.3, was granting us an acceptable error, as shown in the second figure. The error function used to determine the goodness of the estimation is the following:

$$error = ||\hat{x} - x|| \quad (6)$$

in which \hat{x} is the current estimate and x is the true value.

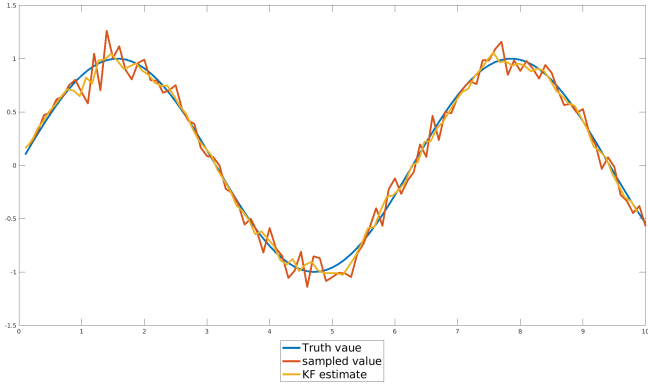


Figure 4: Linear Kalmann Filter test over a sine wave.




Figure 5: Linear Kalmann Filter test over a sine wave error. The red line is the mean.

Once we were satisfied of these results we applied the Kalmann Filter to the data coming from the Lidar. In order to achieve satisfactory results we had to tune up again the observation covariance and process noise matrix, respectively of value 0.3 and 0.1. The tests have been carried out on Lidar data having gaussian noise with 0 mean and variance of 0.01 and 0.1, which resulting point clouds are shown in Figures 6 and 7. The original noisy cloud matrix has been omitted for the clarity of the results.




Figure 6: Estimated point cloud from gaussian noise having 0 mean and 0.01 variance, in red. Noiseless point cloud in blue.




Figure 7: Estimated point cloud from gaussian noise having 0 mean and 0.1 variance, in red. Noiseless point cloud in blue.

As it is possible to notice from the figures, we have that in both cases the Kalmann Filter in use is able to correct the noisy input and produce a result not so different from the truth value, in blue. In particular the process seems to fail in Figure 7 in which a huge area of the ground is shown in red. This error is partially due to the Kalmann Filter itself in estimating the position of the points nearby the ground, but the majority of the fault goes to the SMRF algorithm run to detect the ground which will then be removed from the point cloud. This error also shows how a variance of 0.1 in the input data is able to generate such foggy data to challenge the detection of the ground. Following the needs to have a measure of the goodness of the estimation we tried to compute an estimation error, but due to the misalignment of the estimated points locations with the truth values locations, we had to crop the estimated points resulting into a bayased error plot. Anyway such plot can still offer some hints on the goodness of the estimation. The error plots for both cases are shown in Figures 8 and 9. The error function in use is the following:

$$error = RMSE(||(\hat{x} - x)||) \quad (7)$$

Indeed we can observe that in the first figure the error rises but then saturates, while in the second figure its behaviour results quite jittery.




Figure 8: Estimation error for 0.01 variance. The red line is the mean.




Figure 9: Estimation error for 0.1 variance. The red line is the mean.

B. Pits Detection

As stated in Section III. the pits detection is a core functionality of the solution allowing to avoid unrecoverable states. For this reason from the point clouds it has been extracted the ground for detecting the pits it contains. The criterion used in identifying them is their distance on the Z axis, specifically a point has been classified as a pit one if its location on the Z axis was lower or equal than minus two times the height of the agent. Once the detection ends the points are removed from the ground and added to the occupancy map as occupied space. The results of this approach can be seen in Figures 10 and 11.




Figure 10: Colour mapped point cloud: in gray the natural elements, in green the ground and in red the pits.

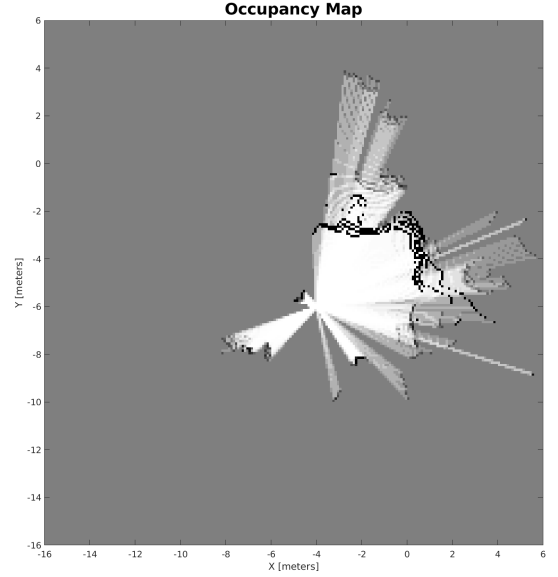


Figure 11: Occupancy map updated with the pits points, upper right.

C. Sharing Data Among Agents

The sharing of data among the agents has been carried out by querying the agentOverviewer class for stating the agents nearby and providing the data of those agents within the communication range to the asking agent. The only issue in this process is the merging and alignment of the data contained in the occupancy maps of the two agents. Thanks to the usage of the absolute coordinates of each agent in simulation and to the role of the agentOverviewer in providing also this information, the resulting merged map is immediately done. Following are the results before and after merging of two agents in opposite locations of the map and having also different orientations. We can notice how in both cases the resulting map is not affected by the different pose of the agents thanks to a preventive homogeneous transformation over the point cloud data into global coordinates performed autonomously by each agent before sharing their data. Instead the transformation of the data in the occupancy map is performed after the sharing by the receiving agent. In Figures 12 and 13 we can observe the before and after merging in the occupancy cloud for Agent 2, while in 14 we can see the resulting merged cloud point.

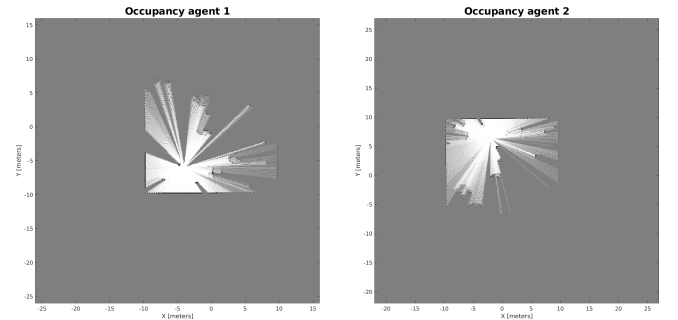


Figure 12: Single occupancy map for the agents in simulation.

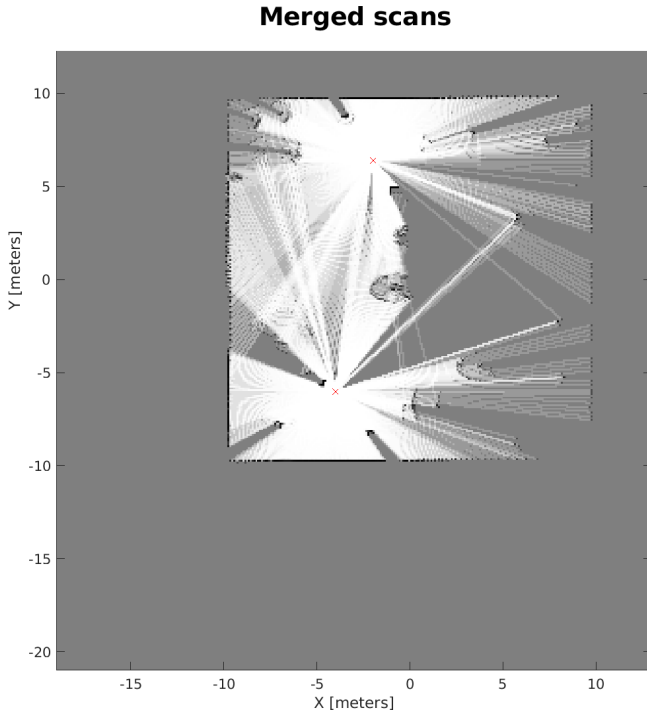


Figure 13: Merged occupancy map for Agent 2. The red crosses are the absolute positions of the agents from which the data were acquired.




Figure 14: Merged point cloud for Agent 2.

D. Roadmap Generation

As stated in Section III. the roadmap computation is carried out by randomly sampling the next location in which to perform the scanning of the environment and computing the path to this position exploiting the RTT* algorithm. The result of this process can be observed in Figure 15. In it we can see that the occupancy map has been inflated in order to avoid collisions with the nearby objects and on top of that the tree to the goal location has been computed.




Figure 15: Computed roadmap, in blue, and selected path to the next location red. Starting point is in the center, arrival point is on top.

E. Agent Location without GPS

The last issue tackled in Section III. is the one of estimating the current absolute pose of an agent employing the Extended Kalmann Filter. In the creation of this filter we used as measurement function the input itself, while for the state transition function we adopted the Differential Drive kinematic model, lastly for the filter implementation we trusted the Matlab implementation. The results in using these settings are as shown in Figures 16, 17 and 18. The tests have been carried out on a custom world in which to the agent was required to follow a path made up of some waypoints, the orange triangles, with the ending location in the upper left corner.




Figure 16: Resulting trajectory of the agent without using EKF.




Figure 17: Resulting trajectory using EKF, red dotted, in a not GPS covered area, green polygon, with the estimated locations with the GPS, blue crosses.

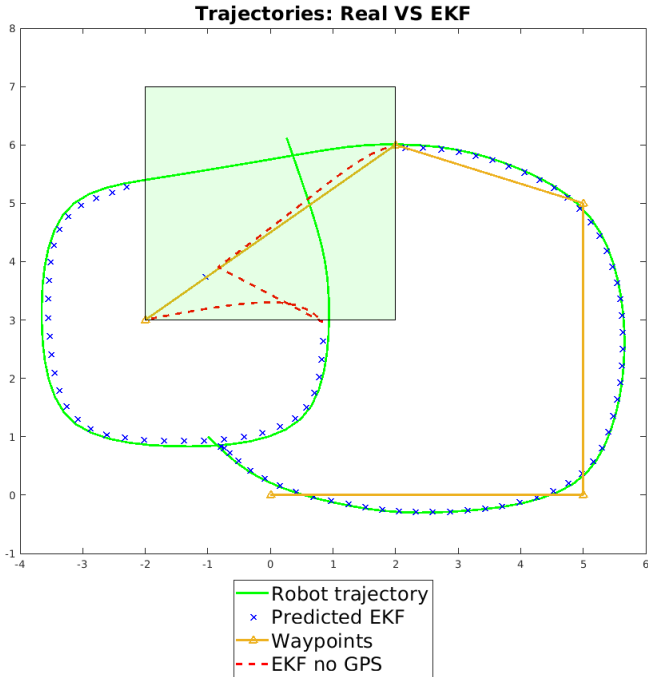


Figure 18: Error in using the EKF with more waypoints.

As we can observe the agent without using the Extended Kalmann Filter carries out its task flawlessly, while using it results into an overshooting of the target due to the mismatch between the estimated location and the real one. The first idea is that the filter has fitted only a straight line which departs from the orientation of the robot when it enters the area without GPS. We firstly supposed that this behaviour was due to a poor tuning of the filter parameters, especially of the process noise matrix which has an important role in fitting non-linearities, but after several test the results were not getting any better. In the latter image it is possible to notice how the filter fails multiple times in reaching the goal point even after the corrections performed outside the area without GPS.

In this report we have presented a solution to perform the Distributed SLAM on a natural environment exploiting terrestrial agents which uses as onboard sensors only a 3D Lidar, a gyroscope, an accelerometer and a proximity sensor. The agents are allowed to share data whenever they are within a communication range between each others. The data from the Lidar have been corrected using a Linear Kalmann Filter, while the location of the agent have been estimated using an Extended Kalmann Filter whenever a reliable GPS location was not available. Our approach provided satisfactory results in the estimation of the data but performed poorly in estimating the absolute pose with the extended filter.

A. Benefits and Limits

Our approach resulted proficient even if the amount of sensors onboard is reduced, but this may result in cheaper applications where the highest amount of the expense is due to the lidar sensor. Unfortunately its inability in estimating correctly its position in absence of GPS data consists in a huge flaw which can lead to unrecoverable states during the exploration. Furthermore, the system is subjected to blind spots in the Lidar point cloud due to the pits in the terrain as shown in Figure 10 being the agent nerby the edge of a cliff. This issue may be easily solved by moving the Lidar sensor in a slightly inclined position facing the ground or even adding a second lidar with a restricted FoV, lastly it is possible to act passively on the problem by reducing the movement distance for the robot and forcing it to acquire more data in despite of the memory requirements and time for the mapping.

B. Future Directions

In the future it should be taken into consideration the presence of water basins, ponds or any other elements containing water in the surroundings, which may be seen as flat surfaces by the laser inside the Lidar, due to the highly reflectance of the water. Knowing this, it should be possible to detect the threat represented by the water through an analysis the intensity of the beam returned to the sensor. Another expansion for this project could be the definition of more advanced classification algorithms to detect the points representing the ground, maybe using SVM or other machine learning techniques for classification. In the end a fix to the Extended Kalmann Filter estimation or an alternative solution must be explored.

The authors have the following addresses: Italy
 {davide.bertelli-1, jhonny.hueller}@studenti.unitn.it.
 This report is the final document for the course of "Distributed Systems for Measurement and Automation".




Figure 19: Full image of the simulated world with pits.